

An Introduction to Sum-Product Networks

A collection of studies on properties, structure, inference and
learning on Sum-Product Networks

Student: Renato Lui Geh

Supervisor: Denis Deratani Mauá (DCC IME-USP)

University of São Paulo / Universidade de São Paulo (USP)

Institute of Mathematics and Statistics / Instituto de Matemática e Estatística
(IME)

Abstract

This work is a collection of ongoing studies I am working on for my undergraduate research project on automatic learning of Sum-Product Networks. The main objective of this work is logging my study notes on this subject in an instructive and uncomplicated way. Most scientific papers are cluttered with intricate names and require extensive background on the subject in order for the reader to understand what is going on. In this paper we seek to provide an easy reference and introductory reading material to those who intend to work with Sum-Product Networks.

This study is divided into five main sections. We start with an introductory section regarding probabilistic graphical models and why Sum-Product Networks are so interesting. Next we talk about the structure of the model. Thirdly, we analyse some properties and theorems. Fourthly, we look on how to perform exact tractable inference. And finally we take a look at how to perform learning.

Contents

1	Introduction	3
1.1	Motivation	3
1.2	Background	4
1.3	Experiments	5
2	Structure of Sum-Product Networks	7
2.1	Partition function	7
2.2	Definition	7
3	Properties of Sum-Product Networks	8
3.1	Completeness and consistency	8
3.2	Validity	8
4	Inference on Sum-Product Networks	8
4.1	Marginals	8
4.2	Most probable explanation (MEP)	8
5	Learning Sum-Product Networks	8
5.1	Learning the weights	8
5.2	Learning the structure	8
	Appendix A Notation	9
A.1	Letters	9
A.2	Events and evidence	9
A.3	Probabilities	9
A.4	Arrows	9
	Appendix B Mathematical background	10
	References	11

1 Introduction

We assume the reader has already read the notation [Appendix A] and has the mathematical background required [Appendix B] defined in the Appendix.

In this section we show what the usual problems with probabilistic graphical models are and what led to the creation of Sum-Product Networks. Additionally, we show some results from experiments Poon and Domingos performed on the inaugural Sum-Product Network article *Sum-Product Networks: A New Deep Architecture* [PD11].

1.1 Motivation

Probabilistic Graphical Models (PGMs) perform inference through posterior probabilities on the query and evidence. Thus, inference would look roughly like this:

$$P(X|\mathbf{e} = e_1, \dots, e_q)$$

Where X is called the variable query and \mathbf{e} the evidence, that is, the observed instances of the variables.

Using the definition of conditional probability,

$$P(X|\mathbf{e}) = \frac{P(X, \mathbf{e})}{P(\mathbf{e})}$$

We get the following equation:

$$P(X|\mathbf{e}) = \frac{P(X, \mathbf{e})}{P(\mathbf{e})} = \alpha P(X, \mathbf{e}) = \alpha \sum_{\mathbf{y}} P(X, \mathbf{e}, \mathbf{y}) \quad (1)$$

Where \mathbf{y} is a hidden variable. That is, let \mathbf{X} be the complete set of variables. Then $\mathbf{X} = \{X\} \cup \mathbf{E} \cup \mathbf{Y}$, where X is the query, \mathbf{E} is the set of evidence variables and \mathbf{Y} is a set of non-query non-evidence variables. Thus \mathbf{y} is an instance of \mathbf{Y} .

We can see that $P(X, \mathbf{e}, \mathbf{y})$ is actually a subset of the full joint distribution. Since we are summing out the hidden variables, we are actually discarding all the possible values of \mathbf{y} and taking into account all the possibilities where the query given the evidence occur.

Now consider a Bayesian network as the PGM of our choice. We know that Bayesian networks have the property of representing the full joint distribution as a product of conditional probabilities:

$$P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{Par}(X_i)) \quad (2)$$

Where $\text{Par}(X_i)$ are the values of the parents of X_i . From this property we know that we can now compute inference by applying Equation (2) on Equation (1). By doing that we get inference by computing the sum of products of conditional probabilities from the network. This is fundamental to Adnan Darwiche's *network polynomial* [Dar03; Dar09], a concept that is the core of Sum-Product Networks.

We know that we can compute inference by summing out the hidden variables and then multiplying the remaining factors, but this process relies on adding and then multiplying an

exponential number of probabilities. In fact, if we don't take the order of the terms in the summation into account, the complexity reaches $O(np^n)$, where p is the number of possible values a variable may take. If we move the independent terms from the summation the complexity is then $O(p^n)$. This is obviously intractable, and a reason why approximate inference is often the best solution.

Bayesian networks are not the only model that have intractable exact inference. Most PGMs suffer from intractability of inference, and hence intractability of learning. However Domingos and Poon argue that “classes of graphical models where inference is tractable exist [...], but are quite limited in the distributions they can represent compactly.” [PD11].

Sum-Product Networks provide a graphical model where inference is both tractable and exact whilst still being more general than existing tractable models.

1.2 Background

In Adnan Darwiche's *A Differential Approach to Inference in Bayesian Networks* [Dar03] and *Modeling and Reasoning with Bayesian Networks* [Dar09], Darwiche presents a new way of representing full joint distributions through a *network polynomial*. In this subsection we will show what a network polynomial is.

Consider the following Bayesian network:




Figure 1: A Bayesian network $A \rightarrow B$, that is, the variable B depends on A .

		A	B	$\Theta_{B A}$
A	Θ_A	a	\bar{b}	$\theta_{b a} = 0.1$
a	$\theta_a = 0.3$	a	\bar{b}	$\theta_{\bar{b} a} = 0.9$
\bar{a}	$\theta_{\bar{a}} = 0.7$	\bar{a}	b	$\theta_{b \bar{a}} = 0.8$
		\bar{a}	\bar{b}	$\theta_{\bar{b} \bar{a}} = 0.2$

Tables 1 and 2

The two tables above describe the Bayesian network in Figure 1. From these tables we can construct the full joint distribution table by applying the definition of conditional probability we saw in the previous subsection.

A	B	$P(A, B)$
a	b	$\theta_a \theta_{b a}$
a	\bar{b}	$\theta_a \theta_{\bar{b} a}$
\bar{a}	b	$\theta_{\bar{a}} \theta_{b \bar{a}}$
\bar{a}	\bar{b}	$\theta_{\bar{a}} \theta_{\bar{b} \bar{a}}$

Table 3 The full joint distribution of the Bayesian network in Figure 1.

Before we proceed we need to understand the concept of variable indicators and their consistency with their respective variables. We will take Boolean variables for simplicity as example. An indicator of a variable, usually written as $[\cdot]$ has a value of 1 if the variable is true and value 0 otherwise. Since $[X_i]$ is quite cumbersome, we abbreviate $[X_i]$ as x_i and $[\bar{X}_i]$ as \bar{x}_i . See [Appendix A] with regards to conflicting notations. Let us now define the notation of indicators more formally.

Definition. Let $\mathbf{X} = \{X_1, \dots, X_n\}$ be the set of all variables in Bayesian network \mathcal{N} , with each variable X_i having p possible values. Then the indicators of an arbitrary variable X_i are denoted by x_{ij} , $1 \leq j \leq p$ where j is the j -th possible value of the variable. All indicators follow the consistency rule.

For $p = 2$, we can abbreviate x_{i1} and x_{i2} to x_i and \bar{x}_i and consider Boolean values. Let us now define consistency. We will define consistency for Boolean values. The extension to multi-valued variables is not discussed here.

Definition. Let $\mathbf{x} = \{x_1, \dots, x_n, \bar{x}_1, \dots, \bar{x}_n\}$ be the set of indicators of all variables in the set $\mathbf{X} = \{X_1, \dots, X_n\}$, where variable X_i may have values 1 or 0, with x_i representing 1 and \bar{x}_i otherwise. Let $\mathbf{e} = \{e_p, \dots, e_q\}$ be the set of an observed event as evidence where e_j represents an observed value for variable X_j in set $\bar{\mathbf{X}}$. Then the set \mathbf{X} is consistent with \mathbf{e} iff:

- For each variable X_i in set \mathbf{X} :
 - If there exists an observed value e_i in \mathbf{e} , then:
 - * If $e_i = 1$, then $x_i = 1$ and $\bar{x}_i = 0$.
 - * If $e_i = 0$, then $x_i = 0$ and $\bar{x}_i = 1$.
 - If there is no observed value e_i for X_i , then:
 - * $x_i = 1$ and $\bar{x}_i = 1$.

We now introduce Darwiche’s network polynomial. Darwiche, in his article and book [Dar03; Dar09], uses indicators as λ_i and $\bar{\lambda}_i$ instead of x_i and \bar{x}_i . However, they are exactly the same as we have defined here. Since we named our variables A and B earlier, we will use Darwiche’s notation for just this example, since it’s more readable. For other examples we will use our own notation.

Table 3 is the full joint distribution that represents the Bayesian network in Figure 1. Darwiche proposes a compact way to represent such distribution by taking each term in the joint distribution, multiplying the relevant indicators, and then summing all terms into a polynomial function. This function is named the network polynomial (Equation 3) of the Bayesian network.

$$\Phi = \lambda_a \lambda_b \theta_a \theta_{b|a} + \lambda_a \bar{\lambda}_b \theta_a \theta_{\bar{b}|a} + \lambda_{\bar{a}} \lambda_b \theta_{\bar{a}} \theta_{b|\bar{a}} + \lambda_{\bar{a}} \bar{\lambda}_b \theta_{\bar{a}} \theta_{\bar{b}|\bar{a}} \quad (3)$$

To compute the probability of any evidence \mathbf{e} , we compute $\Phi(\mathbf{e})$ such that all indicators are consistent with \mathbf{e} . If we assume that the indicators will always be consistent with the evidence, then $\Phi(\mathbf{e})$ becomes the partition function when $\mathbf{e} = \emptyset$. This is true because, if $\mathbf{e} = \emptyset$, then all indicators must be set to 1. Therefore, the value of $\Phi(\mathbf{e})$ must be the highest possible value the function may take. The partition function normalizes an unnormalized distribution.

Now that we know what a network polynomial is, we may start our study on Sum-Product Networks. The next subsection focuses on some experiments and achievements Poon and Domingos performed on [PD11]. The next section introduces Sum-Product Networks.

1.3 Experiments

In this subsection we take a brief look at some of the results Poon and Domingos worked on on their article *Sum-Product Networks: A New Deep Architecture* [PD11].

Conducting experiments on two sets of image, Caltech-101 and the Olivetti face dataset, Poon and Domingos achieved astounding results. With the lowest mean squared error compared to Deep Boltzmann (DBMs), Deep Belief Networks (DBN), Principal Component Analysis (PCA) and Nearest Neighbour (NN), Sum-Product Networks (SPNs) outperformed all the other models.

Learning Caltech faces took 6 minutes with 20 CPUs, learning for DBMs/DBNs ranged from 30 hours to over a week.[PD11]

“For inference, SPNs took less than a second to find the MPE completion of an image, to compute the likelihood of such a completion, or to compute the marginal probability of a variable [...]. In contrast, estimating likelihood is a very challenging problem; estimating marginals requires many Gibbs sampling steps that may take minutes of even hours, and the results are approximate without guarantee on the quality.” [PD11]




Figure 2: Image completion output from SPNs.

Poon and Domingos also conducted preliminary experiments on object recognition. They ran classification on three classes (one vs the other two) against convolutional DBNs (CDBNs). SPNs showed almost flawless results.

Architecture	Faces	Motorbikes	Cars
SPN	99%	99%	98%
CDBN	95%	81%	87%

Table 4 Comparison between CDBN and SPN on object recognition.

More information can be found on Poon and Domingos *Sum-Product Networks: A New Deep Architecture* [PD11] and *Learning the Structure of Sum-Product Networks* [GD13].

2 Structure of Sum-Product Networks

In this section we work on the ideas presented in the previous sections. We will base this section on the works of Poon and Domingos on their article *Sum-Product Networks: A New Deep Architecture* [PD11] and also on Gens and Domingos' *Learning the Structure of Sum-Product Networks* [GD13].

2.1 Partition function

Probabilistic graphical models represent distributions compactly through a normalized product of factors.

$$P(X = x) = \frac{1}{Z} \prod_k \phi_k(x_{\{k\}})$$

This equation states that the probability of X taking value x is the product of all factors, with each factor taking a subset of the variables, divided by a partition function. The partition function is the result of all the product of factors summed out.

$$Z = \sum_{x \in \mathcal{X}} \prod_k \phi_k(x_{\{k\}})$$

However, this form cannot represent all distributions. Additionally, inference is exponential in scope size in the worst-case and learning takes sample size and time exponential. This is because the partition function Z is the sum of an exponential number of terms, and since all marginals are sums of these terms, if Z is tractable, then marginals are also tractable.

In Domingos and Poon [PD11], they claim that, since “ Z is computed using only two types of operations sums and products”, then “it can be computed efficiently if $\sum_{x \in \mathcal{X}} \prod_k \phi_k(x_{\{k\}})$ can be reorganized using the distributive law into a computation involving only a polynomial number of sums and products.” From this idea, they elaborate on Darwiche’s network polynomial and introduce Sum-Product Networks, a representation that seeks to admit an efficient form for Z whilst still being general.

2.2 Definition

Now we define a sum-product network formally.

Definition. A *sum-product network (SPN)* is an *acyclic digraph*. An SPN S over variables X_1, \dots, X_n has leaves as indicators $x_1, \dots, x_n, \bar{x}_1, \dots, \bar{x}_n$ and internal nodes as sums and nodes in alternating layers. For each edge (i, j) that emanates from a sum node i , there exists a non-zero weight w_{ij} . The value of a sum node i is $\sum_{j \in Ch(i)} w_{ij} v_j$, where $Ch(i)$ is the set of children of i and v_j is the value of the node j . The value of a product node i is $\prod_{j \in Ch(i)} v_j$. The value of an SPN is the value of its root.

An SPN S as a function of the indicator variables $x_1, \dots, x_n, \bar{x}_1, \dots, \bar{x}_n$ will be denoted by $S(x_1, \dots, x_n, \bar{x}_1, \dots, \bar{x}_n)$. A complete state x means that, for each variable X_i , there are two possible outcomes for the indicators: either $x_i = 1, \bar{x}_i = 0$ or $x_i = 0, \bar{x}_i = 1$. When an SPN S has a complete state x , instead of enumerating each indicator, we simply abbreviate to $S(x)$. $S(e)$ will be used to denote when the indicators are specified according to the evidence e .

A sum-product network is an elaboration on Darwiche’s network polynomial. We can look at an SPN as a way to graphically represent the network polynomial of the distribution.




Figure 3: Two sum-product networks. The one on the left implements a naive Bayes mixture model and the other implements a junction tree.

The partition function is the value of the network polynomial when all indicators are set to 1. That is, $Z = \Phi(\mathbf{e})$ when $\mathbf{e} = \emptyset$. If an SPN S has all indicators set to 1, we refer to it as $S(*)$. As we will see later, if an SPN S is valid, then $S(*) = Z_S$.

The probability of evidence, $P(e) = \Phi(e)/Z$, is linear in the size of the network polynomial, which in turn is exponential in the number of variables. However, it is possible, with a sum-product network, to represent and evaluate the probability of evidence in space and time polynomial.

3 Properties of Sum-Product Networks

3.1 Completeness and consistency

3.2 Validity

4 Inference on Sum-Product Networks

4.1 Marginals

4.2 Most probable explanation (MEP)

5 Learning Sum-Product Networks

5.1 Learning the weights

5.2 Learning the structure

A Notation

In this section we show the notations we use throughout this paper.

A.1 Letters

We use an uppercase letter to denote a variable. A lowercase letter denotes an instance of a variable. A bold fonted letter is a set. A bold fonted uppercase letter is a set of variables. For instance:

$$\mathbf{X} = \{X_1 = x_1, \dots, X_n = x_n\}$$

When dealing with indicator variables, the indicator function will be abbreviated as a lower-case letter. That is, $[X_i]$ will be abbreviated to x_i . Similarly, $[\overline{X}_i]$ will be written as \overline{x}_i . This clearly contradicts our first notation rule, however the two meanings will be clear from context and therefore conflicts will never occur.

A.2 Events and evidence

The letter ‘e’, regardless of case, is reserved for events and evidence. An uppercase ‘ E ’ is an evidence variable. An uppercase bold fonted ‘ \mathbf{E} ’ is the set of evidence variables. A lowercase ‘ e ’ is a particular observed event variable. A bold fonted lowercase ‘ \mathbf{e} ’ is the set of variables of a particular observed event.

A.3 Probabilities

All functions of the form $P(\cdot)$ are probability functions. Joint probability distributions have the variables separated by commas $P(X, Y)$ instead of $P(X \wedge Y)$. We call prior probabilities the probability functions of the form $P(X)$. Posterior probabilities are of the form $P(\mathbf{X}|\mathbf{Y})$.

When enumerating a set of instances we may omit commas, brackets or set name. For instance:

$$\begin{aligned} P(\mathbf{X} = \overline{a}bc) &\text{ is equivalent to } \\ P(\mathbf{X} = \{a, \overline{b}, c\}) &\text{ is equivalent to } \\ P(\overline{a}bc) &\text{ is equivalent to } \\ P(a = \text{true}, b = \text{false}, c = \text{true}) \end{aligned}$$

A.4 Arrows

An arrow pointing to the right may have two possible meanings:

- Dependency
 - $A \rightarrow B$ is read as *B depends on A*.
- Directed edge connectivity
 - $A \rightarrow B$ is read as *there exists an edge from node A to node B*.

Meanings will be clear from context.

B Mathematical background

References

- [Dar03] Adnan Darwiche. “A Differential Approach to Inference in Bayesian Networks”. In: (2003).
- [Dar09] Adnan Darwiche. *Modeling and Reasoning with Bayesian Networks*. 1st Edition. Cambridge University Press, 2009.
- [GD13] Robert Gens and Pedro Domingos. “Learning the Structure of Sum-Product Networks”. In: *International Conference on Machine Learning* 30 (2013).
- [PD11] Hoifung Poon and Pedro Domingos. “Sum-Product Networks: A New Deep Architecture”. In: *Uncertainty in Artificial Intelligence* 27 (2011).